

# Statistical Neuroscience in the Single Trial Limit

Alex H. Williams and Scott W. Linderman

March 10, 2021

## Abstract

Individual neurons often produce highly variable responses over nominally identical trials, reflecting a mixture of intrinsic “noise” and systematic changes in the animal’s cognitive and behavioral state. In addition to investigating how noise and state changes impact neural computation, statistical models of trial-to-trial variability are becoming increasingly important as experimentalists aspire to study naturalistic animal behaviors, which never repeat themselves exactly and may rarely do so even approximately. Estimating the basic features of neural response distributions may seem impossible in this trial-limited regime. Fortunately, by identifying and leveraging simplifying structure in neural data—e.g. shared gain modulations across neural subpopulations, temporal smoothness in neural firing rates, and correlations in responses across behavioral conditions—statistical estimation often remains tractable in practice. We review recent advances in statistical neuroscience that illustrate this trend and have enabled novel insights into the trial-by-trial operation of neural circuits.

## Introduction

Widely disseminated optical and electrophysiological recording technologies now enable many research labs to simultaneously record from hundreds, if not thousands, of neurons. Often, a first step towards characterizing the resulting datasets is to estimate the *average* response of all neurons across a small set of *conditions*. For example, a subject may be presented different sensory stimuli (images, odors, sounds, etc.) or trained to perform different behaviors (reaching to a target, pressing a lever, etc.), each of which constitutes a different condition. Each condition is then repeated many times, and the neural response is averaged over these nominally identical *trials* to reduce noise and variability.

Despite their obvious importance, averages represent incomplete (and potentially misleading [28]) summaries of neural data. Trial-to-trial variations in neural activity reflect a variety of interesting processes, including fluctuations in attention and task engagement [12<sup>•</sup>, 68<sup>•</sup>, 73, 84], changes-of-mind during decision-making [13, 23, 33, 34, 69], modulations of behavioral variability to promote learning [14], representations of uncertainty [49], changes-in-strategy [72], and modified sensory processing linked to active sensing [22, 91], locomotion [94], and other motor movements [53, 88]. Some of these effects may be disentangled by developing targeted experimental designs [84] or by analyzing behavioral covariates [53, 55, 88]. In other cases, these effects may spontaneously emerge and subside during the course of an experiment and leave little to no behavioral


signature.

While the activity of individual neurons may correlate with these single-trial phenomena, such effects may be subtle and difficult to detect. A unique advantage to collecting *simultaneous* population recordings is the possibility of pooling statistical power across many individually noisy neurons to characterize short-term fluctuations and long-term drifts in neural circuit activity. One way to approach this goal is to estimate higher-order statistics of the neural response distribution, such as the trial-to-trial response correlations between all pairs of neurons. Unfortunately, as we discuss below, estimating second-order response properties (i.e., covariance structure) generally requires the number of trials per condition to grow super-linearly with the number of neurons, which can quickly become infeasible. However, we will see that there are good reasons to believe this worst-case analysis is overly pessimistic, and that it can be overcome by carefully designed statistical analyses.

At the same time, there is a trend toward studying neural circuits in more ethologically relevant settings. This involves studying rich sensory stimuli and spontaneous, unconstrained behaviors which elicit more natural patterns of neural activity. Recent work in visual neuroscience, for instance, has measured activity in response to very diverse sets of natural images [7, 95], with as little as two trials per image [87]. This starkly contrasts with classical experiments, which presented simple stimuli (e.g. oriented gratings) repeatedly over many trials. Similar trends are present in motor neuroscience, where motion capture algorithms have been leveraged to measure spontaneous animal behaviors [46, 47]. Unconstrained motor actions repeat themselves infrequently and inexactly, resulting in few “trials” compared with classical behavioral tasks (e.g. cued point-to-point reaches or lever presses).

Figure 1A summarizes these trends. We selected a small subset of papers from the past thirty years that obtained multi-neuronal recordings, and plotted the total number of trials collected against the number of free variables one could potentially try to estimate. There are a total of  $NC$  such variables for a recording of  $N$  neurons across  $C$  conditions. (For now, we neglect the within-trial temporal dynamics of neural responses; including these dynamics as estimatable parameters would only exacerbate the potential for statistical error.) The greyscale background shows the expected estimation error for second-order statistics under a worst-case scenario where trial-to-trial variations are decorrelated across neurons and conditions (i.e., variability that is high-dimensional, in a sense that we formally define below). We observe that the number of trials has been growing more slowly than number of parameters we wish to estimate, raising the possibility that our statistical analysis will suffer from greater estimation errors. If these trends continue, traditional experimental designs that assume large numbers of trials over a discrete set conditions, will become an increasingly ill-suited framework for neural data analysis—under completely unconstrained and naturalistic settings, no two experiences and actions are truly identical so, in some sense, each constitutes a unique condition with exactly one trial.

The solution to this apparent problem is to recognize that neural and behavioral data may not be as high dimensional as they appear. Though we may never see the exact same pattern of neural activity or postural dynamics twice, our measurements may lie close to a low dimensional manifold. For example, the neuron-by-neuron covariance matrix might be approximately low rank, trial-to-trial variability may arise from a small



**Figure 1:** (A) The number of trials in neural datasets is growing more slowly the number of simultaneously recorded neurons and sampled behavioral conditions. Scatter plot color corresponds to year of publication on an ordinal scale (see legend). Grayscale heatmap shows the worst-case error scaling for covariance estimation [92°]—the contours are  $O(NC \log NC)$  for a dataset with  $N$  neurons and  $C$  conditions. Darker shades correspond to larger error. (B) Low-dimensional visualizations of trial-to-trial variability in static (top row) and dynamic (bottom row) neural responses. *Left*, trial average in two conditions (blue and red). In the dynamic setting, neural firing rates evolve along a 1D curve parameterized by time. In the static setting, responses are isolated points in firing rate space. *Middle*, same responses but with independent single-trial variability illustrated in each dimension. *Right*, same responses with correlated variability. The positive correlations in the top panel are “information limiting” because they increase the overlap between the two response distributions, degrading the discriminability of the two conditions (see, e.g., [2]). In the bottom panel, correlations in neural response amplitudes result in trajectories that are preferentially stretched or compressed along particular dimensions from trial-to-trial (see [102°] for a class of models that are adapted to this simplifying structure).

number of internal state variables, and natural behavior may be composed of a small number of relatively stereotyped movements. If and when such simplifying structure is found, it can dramatically reduce the number of trials necessary to obtain accurate parameter estimates. The success of single trial analysis of neural and behavioral data relies crucially on our ability to identify and leverage these patterns in our data. The rest of this review summarizes recent examples from the neuroscience literature that exemplify this approach.

## Statistical challenges in trial-limited regimes

In a pioneering study from the early 1990’s, Zohary et al. [109] measured responses from co-recorded neuron pairs over 100 sessions in primates performing a visual discrimination task. They found weak, but detectable, correlations in the neural responses over trials—when one neuron responded with a large number of spikes, the co-recorded neuron often had a slightly higher probability of emitting a large spike count. Though this result may appear innocuous, the authors were quick to point out that even weak correlations could drastically impact the signalling capacity of sensory cortex (fig. 1B). This finding inspired a large number of experimental and theoretical investigations into “noise correlations,” which today represents one of the most developed bodies of scientific work on trial-to-trial variability (for reviews, see [2, 36]).

The fact that Zohary et al. [109] were technologically limited to recording two neurons at once came with a silver lining. For every variable of interest (i.e. a correlation coefficient between a unique pair of neurons),

### Box 1 — Notation

We aimed to keep mathematical notation light, but we summarize a few points of standard notation here. We denote scalar variables with non-boldface letters (e.g.  $x$  or  $X$ ), vectors with lowercase, boldface letters (e.g.  $\mathbf{x}$ ), and matrices with uppercase, boldface letters (e.g.,  $\mathbf{X}$ ). The set of real numbers is denoted by  $\mathbb{R}$ , so the expression  $s \in \mathbb{R}$  means that  $s$  is a scalar variable. Likewise,  $\mathbf{v} \in \mathbb{R}^n$  means that  $\mathbf{v}$  is a length- $n$  vector, and  $\mathbf{M} \in \mathbb{R}^{m \times n}$  means that  $\mathbf{M}$  is a  $m \times n$  matrix. A matrix  $\mathbf{X} \in \mathbb{R}^{m \times n}$  is said to be “low-rank” if there exist matrices  $\mathbf{U} \in \mathbb{R}^{m \times r}$  and  $\mathbf{V} \in \mathbb{R}^{n \times r}$ , where  $r < \min(m, n)$  and for which  $\mathbf{X} = \mathbf{U}\mathbf{V}^\top$ . The smallest value of  $r$  for which this is possible is called the *rank* of  $\mathbf{X}$ , in which case we would say “ $\mathbf{X}$  is a rank- $r$  matrix.” We briefly utilize “*Big O notation*” to represent a function up to a positive scaling constant. More precisely, we will use  $O(g(N))$  to represent an anonymous function that is upper bounded by  $g(N)$ , up to an absolute constant. That is,  $O(g(N))$  represents any function  $f$  that satisfies  $f(N) \leq C \cdot g(N)$  for some constant  $C > 0$ . When  $g$  is a monotonically increasing function of  $N$ , dropping constant terms like this is useful to understand the scaling behavior of the system in the limit as  $N$  becomes very large.

they collected a large number of independent trials. Their statistical analyses were straightforward because the number of unknown variables was much smaller than the number of observations. Two recent studies by Bartolo et al. [3] and Rumyantsev et al. [74<sup>•</sup>] revisited this question using modern experimental techniques. The latter group recorded calcium-gated fluorescence traces from  $N \approx 1000$  neurons in mouse visual cortex over  $K \approx 600$  trials. Since each session contained roughly  $N(N - 1)/2 = 499500$  pairs of neurons, the number of free parameters (correlation coefficients between unique neuron pairs) was vastly larger than the number of independent observations. Intuitively, if the correlations between neuron pair A-B and neuron pair B-C were mis-estimated, then the estimated correlation between neurons A and C would also likely be inaccurate, since the same set of trials were used for the underlying calculation. The authors were forced to grapple with an increasingly common question: how many neurons and trials must be collected to ensure the overall conclusions were accurate?

The field of high-dimensional statistics provides answers to questions like this [92<sup>•</sup>, 97]. A standard result states that, in order to accurately estimate a  $N \times N$  covariance matrix,  $\Sigma$ , the number of trials should be  $O(d \log N)$  where  $d$  measures the *effective dimensionality* of the covariance matrix. If the tails of the neural response distribution decay sufficiently fast, this bound can be improved to  $O(d)$  trials [93]. Formally, the dimensionality is defined as  $d = \text{Tr}[\Sigma]/\|\Sigma\|$  (see section 5.6 of [92<sup>•</sup>]). Intuitively,  $d$  is large when trial-to-trial variability equally explores every dimension of neural firing rate space (as in fig. 1B, “uncorrelated variability”). Conversely,  $d$  is small when there are large correlations in a small number of dimensions, such that many dimensions are hardly explored relative to high-variance dimensions. In the worst case scenario where variability is high-dimensional and heavy-tailed, we would have  $d = N$  and require  $O(N \log N)$  trials, which may be infeasible to collect.

Fortunately, Rumyantsev et al. [74<sup>•</sup>] provide empirical evidence and a simple circuit model which suggest that the eigenvalues of  $\Sigma$  decay rapidly. This corresponds to the statistically tractable setting where  $d$  is much smaller than  $N$ . Bartolo et al. [3] contemporaneously reported similar results under different experimental conditions and in nonhuman primates. Overall, these works provide some of the strongest evidence to date

that noise correlations do indeed limit the information content of large neural populations. The effect of noise correlations more generally—e.g., in concert with behavioral state changes [54], and in response to more diverse stimuli [71]—remains a subject of active research.

For us, the key takeaway is that the presence of simplifying structures (e.g. low dimensionality) enables accurate statistical analysis when we collect data from many more neurons and behavioral conditions than trials. We can leverage formal results from high-dimensional statistics to sharpen this conceptual lesson into quantitative guidance for experimental designs.


## Gain modulation and low-rank matrix decomposition

It is useful to view covariance estimation (discussed above), as a special case of *maximum likelihood estimation* (MLE). Given a recording of  $N$  neurons, we observe neural responses  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_K\}$ , where each  $\mathbf{x}_k \in \mathbb{R}^N$  is a vector holding the population response on trial  $k \in \{1, 2, \dots, K\}$ . Now, assume that each response is sampled independently from a multivariate normal distribution with mean  $\mu$  and covariance  $\Sigma$ ; that is,  $\mathbf{x}_k \sim \mathcal{N}(\mu, \Sigma)$  for every trial index  $k$ . It is a simple exercise to show that, under a log-likelihood objective function, the maximum likelihood parameter estimates are the empirical mean,  $\hat{\mu} = \frac{1}{K} \sum_{k=1}^K \mathbf{x}_k$ , and the empirical covariance,  $\hat{\Sigma} = \frac{1}{K} \sum_{k=1}^K (\mathbf{x}_k - \hat{\mu})(\mathbf{x}_k - \hat{\mu})^\top$ .

In the last section, we discussed a scenario where the covariance was well-fit by a low-rank decomposition—i.e., there is some  $N \times d$  matrix  $U$ , for which  $\Sigma \approx UU^\top$ —and saw that such structure allows us to estimate the covariance matrix accurately in a reasonable number of trials. Other modeling assumptions can achieve similar effects. For example, Wu et al. [105<sup>•</sup>] developed a model with *Kronecker product* structure to model variability across multiple data modalities (odor conditions and neurons). Integrating data from multiple modalities into a unified model is a nascent theme of recent research in neuroscience [20, 51, 59, 80, 102<sup>•</sup>], and is already well-established in other areas of computational biology [108<sup>•</sup>].

Using a multivariate normal distribution to model single-trial responses is often mathematically convenient and computationally expedient. However, it is usually not an ideal model. For example, if we count the number of spikes in small time windows as a measure of neural activity, a Poisson distribution is typically used to model variability at the level of single neurons [61]. Extending the Poisson distribution to the multivariate setting turns out to be a somewhat advanced and nuanced subject [31]. A simple approach is to introduce per-trial *latent variables*, which induce correlated fluctuations across neurons. For example, suppose the number of spikes fired by neuron  $n$  on trial  $k$  is modeled as  $X_{nk} \sim \text{Poisson}(\mathbf{u}_n^\top \mathbf{v}_k)$ , where  $\mathbf{u}_n \in \mathbb{R}^r$  and  $\mathbf{v}_k \in \mathbb{R}^r$  are vectors holding  $r$  “components” or latent variables for each neuron and trial. It is useful to reformulate this model using matrix notation. Let  $\Lambda = UV^\top$ , where the matrices  $U \in \mathbb{R}^{N \times r}$  and  $V \in \mathbb{R}^{K \times r}$  are constructed by stacking the vectors  $\mathbf{u}_n$  and  $\mathbf{v}_k$ , row-wise. We can interpret  $\Lambda$  as a matrix of estimated firing rates, which, in expectation, equals the observed spikes counts—i.e., our model is that  $\mathbb{E}[X] = \Lambda = UV^\top$ , where  $X$  is an  $N \times K$  matrix holding the observed spike counts, and noise is Poisson-distributed (i.i.d. across neurons and trials).

This model is a special case of a *low-rank matrix factorization* (Box 2)—a versatile framework that encompasses



**Figure 2:** Matrix factorization methods for single-trial analysis. **(A)** Schematic firing rate traces of three neurons demonstrating correlated gain-modulation: the peak responses in all three neurons are scaled by a common factor on each trial. **(B)** A rank-1 NMF model over 9 neurons and 9 trials. The neural responses on each trial are taken to be the peak evoked firing rate as illustrated in panel A. The data  $X$  (black-to-white heatmap) are approximated by the outer product of two vectors,  $uv^\top$  (respectively shown as blue and red stem plots). **(C)** Interpretation of a rank-3 NMF model as an idealized neural circuit. Low-dimensional neuron factors,  $U$ , correspond to synaptic weights, while trial factors,  $v_k$  for trial  $k$ , correspond to input amplitudes (a related circuit interpretation is provided in [102\*]). **(D)** A schematic data matrix containing responses from 9 neurons over 9 trials is modeled as the sum of three rank-1 components (overall, a rank-3 model). **(E)** Spike-triggered ensemble analysis of retinal ganglion neurons in Salamander retina. NMF-identified components correspond to localized visual inputs that correspond to presynaptic bipolar cells; these signals are mixed together in PCA-identified components (panel adapted from [44]). **(F)** Demonstration of the “rotation problem” in PCA. *Left*, a rank-three data matrix holding a multivariate time series. *Right*, temporal factors identified by PCA and NMF (colored lines). Dashed black line denotes zero loading. In this case, the decomposition by NMF is unique (up to permutations; [17]), unlike PCA.

many familiar methods including PCA, k-means clustering, and others [90<sup>•</sup>]. Now we ask: given  $X$ , what are good values for  $U$  and  $V$ ? One way to answer this question is to optimize  $U$  and  $V$  with respect to the log-likelihood of the data. After removing an additive constant from the log-likelihood function, we arrive at the following optimization problem:

$$\begin{aligned} \text{maximize} \quad & \sum_{n=1}^N \sum_{k=1}^K X_{nk} \log \Lambda_{nk} - \Lambda_{nk} \\ \text{subject to} \quad & \Lambda = UV^\top; \quad U \geq 0; \quad V \geq 0. \end{aligned} \tag{1}$$

We have included nonnegativity constraints which ensure that  $\Lambda \geq 0$ ; i.e., the predicted firing rates are, quite sensibly, nonnegative. Alternatively, we could have dropped the nonnegativity constraints, and set  $\Lambda = \exp(UV^\top)$ , where the exponential is applied elementwise. However, the factor matrices are typically easier to interpret when they are nonnegative, as demonstrated by Lee and Seung [38], who popularized the model in eq. (1) under the name of *nonnegative matrix factorization* (NMF). An alternative form of NMF is nearly identical, but uses a least-squares criterion instead of the Poisson log-likelihood (see [25] for a modern review).

NMF is of immense importance to modern neural data analysis. Variants of this method underlie popular algorithms for spike sorting [85], extracting calcium fluorescence traces from video [66, 107], identifying sequential firing patterns in neural time series [45, 65], parcellating widefield imaging videos into functional regions [78], and theories of grid cell pattern formation [18, 83]. In the context of single-trial analysis, NMF can be interpreted as a model of *gain modulation* (fig. 2A), which is a widely studied phenomenon in sensory and motor circuits [21, 63]. For example, in a rank-1 NMF model, the firing rate matrix is factorized by a pair of vectors,  $\Lambda = u v^\top$  (fig. 2B). We can interpret  $u$  as being proportional to the trial-averaged firing rate of all  $N$  neurons. On trial  $k$ , the predicted firing rates are  $\Lambda(:, k) = v_k u$ , which is simply the average response re-scaled by a per-trial gain factor,  $v_k$ . It has been hypothesized that such gain modulations play a key role in tuning the signal-to-noise ratio of sensory representations, with larger gain factors corresponding to attended inputs [68<sup>•</sup>, 70].

NMF can also model more complex patterns of single-trial variability. In particular, an NMF model with  $r$  low-dimensional components (i.e. a rank- $r$  model) can capture independent gain modulations over  $r$  neural sub-populations (fig. 2C-D). Despite differing in some important details, recent statistical models of sensory cortex bear some similarity to this framework [68<sup>•</sup>, 98, 99]. We thus view this conceptual connection between NMF—a general-purpose method with many applications outside of neuroscience [25]—and the neurobiological principle of gain modulation as a useful and unifying intuition.

Recent work has also applied NMF to *spike-triggered analysis* of visually-responsive neurons. Here, the data matrix  $X$  is a  $S \times K$  matrix holding the spike-triggered ensemble: the  $k^{\text{th}}$  column of  $X$  contains the visual stimulus (reshaped into a vector) that evoked the  $k^{\text{th}}$  spike in the recorded neuron. In this context, trial-to-trial variability corresponds to variation in the stimulus preceding each spike. The classic method of spike-triggered covariance

## Box 2 — Matrix Factorization

The expression  $\widehat{X} = UV^\top$  is called a *matrix factorization* (or *matrix decomposition*) of  $\widehat{X}$ . We call  $U$  and  $V$  *factor matrices* and the columns of these matrices *factors*. This terminology is analogous to factoring natural numbers: the expression  $112 = 7 \cdot 16$  is a factorization of 112 into the factors 7 and 16. Low-rank factorizations are a common element of many statistical models: given a data matrix  $X \in \mathbb{R}^{m \times n}$ , these models posit a low-rank approximation  $\widehat{X} = UV^\top$ , where  $U \in \mathbb{R}^{m \times r}$ ,  $V \in \mathbb{R}^{n \times r}$  and  $r < \min(m, n)$  is the rank of  $\widehat{X}$ . This model summarizes the  $mn$  datapoints held in  $X$  with  $mr + nr$  model parameters, which may be substantially smaller. Intuitively, this corresponds to a form of *dimensionality reduction* since the dataset is compressed into an low-dimensional (i.e.,  $r$ -dimensional) subspace.


Adding constraints and regularization terms on the factor matrices is often very useful [90<sup>•</sup>]. For example, if  $F \in \mathbb{R}^{r \times r}$  denotes an arbitrary invertible matrix, and if  $U$  and  $V$  are optimal factor matrices by a log-likelihood criterion, then  $UF^{-1}$  and  $VF^\top$  are also optimal factor matrices since  $(UF^{-1})(VF^\top)^\top = UF^{-1}FV^\top = UV^\top$ . In the context of PCA and factor analysis, this invariance is called the “rotation problem.” This degeneracy of solutions hinders the interpretability of the model—ideally, the pair of factor matrices that maximize the log likelihood would be unique (up to permutation and re-scaling components). This can sometimes be accomplished by adding nonnegativity constraints (as in NMF; [17]) or L1 regularization (as in sparse PCA; [112]). The tensor factorization model we discuss in this review (see [102<sup>•</sup>]) are also essentially unique under mild conditions. Kolda and Bader [37] review these conditions and other forms of tensor factorizations (namely, Tucker decompositions) that do not yield unique factors.

analysis [79], captures this variability by a low-rank decomposition of the empirical covariance matrix—in essence, applying PCA to  $X$ . Liu et al. [44] showed that NMF extracts more interesting and physiologically interpretable structure. When applied to data from retinal ganglion cells, NMF factors closely matched the location of presynaptic bipolar cell receptive fields, which were identified by independent experimental measurements (fig. 2E). Subsequent work by Shah et al. [81<sup>•</sup>] sharpened the spike-triggered NMF model in several respects to achieve impressive results on nonhuman primate retinal cells and V1 neurons.

The one-to-one matching of NMF factors to interpretable real-world quantities is a remarkable capability. In contrast, the low-dimensional factors derived from PCA generally *cannot* be interpreted in this manner due to the “rotation problem” described in **Box 2** and illustrated in Figure 2F. In essence, PCA only identifies the linear subspace containing maximal variance in the data, and there are multiple equivalent coordinate systems that describe this subspace. Under certain conditions (outlined in [17]), the additional constraints in the NMF objective cause the solution to be “essentially unique” (that is, unique up to permutations and scaling transformations). This markedly facilitates our ability to interpret the features derived from NMF, and is one of the main explanations for the method’s widespread success.

## Single-trial variability in temporal dynamics

Thus far, we have discussed models that treat neural responses as a static quantity on each trial. However, theories spanning motor control [96], decision-making [27], odor discrimination [104], and many other



**Figure 3:** Models for trial-to-trial variability in temporal dynamics (A) Left, correlation between pursuit latency (reaction time) and neural response latency in a single-neuron recording from a nonhuman primate performing smooth eye pursuits. Right, correlation in response latency between two co-recorded neurons in the same task. All units are z-scored. Large black dots denote averages over quintiles. (Adapted from [39].) (B) Left, diagram illustrating three trials containing four behavioral actions. Right, time warping functions that align the three trials. The “canonical trial” defines the identity line, and time is re-scaled on the two other trials to align each behavioral action. This manual alignment procedure is used in refs. [35, 40]. (C) Activity from a neuron in rat motor cortex with spike times aligned to lever press (red arrow, left) and aligned by unsupervised time warping (right). Importantly, the warping functions were fit only to neural data from simultaneously recorded neurons—the discovery of spike time oscillations demonstrates that variability in timing is correlated across cells, and thus generalizes effectively on this heldout neuron. (Adapted from [103\*].) (D) Illustration of how trial-to-trial variability in time warping and gain-modulated population dynamics respectively affect the speed and scale of firing rate trajectories. Top panels show firing rate traces from two neurons, while bottom plots show the trajectory in a low-dimensional state space. (E) Schematic illustration of the canonical tensor decomposition model. (F) A set of three low-dimensional factors derived from tensor decomposition. The model identifies a sub-population of neurons (blue) whose within-trial dynamics (red) grow in amplitude at the onset of learning (green). (Adapted from [102\*].)

domains, all predict that the *temporal* dynamics of neural circuits are crucial determinants of behavior and cognition. Similar to how trial-to-trial variability in neural amplitudes can be described by a small number of shared gain factors, variability in temporal activity patterns also tends to be shared across neurons. For example, the latency of neural responses can correlate with behavioral reaction times on a trial-by-trial basis; similarly, the latencies of co-recorded neuron pairs are often correlated [1, 39] (fig. 3A).

The temporal patterning of population dynamics can also vary in more complex ways. In many experiments, each trial is composed of a sequence of sensory cues, behavioral actions, and reward dispensations. The time delays between successive events often vary on a trial-by-trial basis, resulting in *nonlinear* warping in the time course of neural dynamics (fig. 3B). These misalignments can obscure salient features of the neural dynamics, and thus is it often crucial to correct for them. This can be done in a human-supervised fashion by warping the time axis to align salient sensory and behavioral key points across trials [35, 40]. Recent work has demonstrated the effectiveness of *unsupervised* time warping models, which are fit purely on neural population data and are agnostic to alignment points identified by human experts [19, 103<sup>•</sup>]. Such models provide a data-driven approach for time series alignment, which can uncover unexpected features in the data—for example, Williams et al. [103<sup>•</sup>] found ~7 Hz oscillations in spike trains from rodent motor cortex, which were not time-locked to paw movements or the local field potential (fig. 3C).

Beyond time warping, we are also interested in trial-to-trial changes in the trajectory's *shape*. For example, geometrical features like rotations [10], tangling [75], divergence [76], and curvature [82<sup>•</sup>], have been used to describe population dynamics in primary and supplementary motor cortex. While such analyses are often carried out at the level of trial averages, estimates of single-trial trajectories can be derived via well-established methods including PCA [11], Gaussian Process Factor Analysis (GPFA; [106]), and, more recently, artificial neural networks [60].

However, some datasets contain thousands of trials from the same population of neurons over multiple days [16]. In these cases, applying methods like PCA and GPFA produces thousands of estimated low-dimensional trajectories (one for each trial). It is infeasible to visually digest and interpret such a large number of trajectories when they are overlaid on the same plot. However, if we hypothesize that trial-to-trial variation is well-described by simple transformations—e.g., time warping or gain modulation (fig. 3D)—we can develop statistical models that exploit this structure to reduce the dimensionality *both* across trials and neurons.

*Tensor factorization* (or tensor decomposition) models [37] can be used to model population activity as a small number of temporal factors that are gain-modulated across trials. Neural data are represented in a 3-dimensional data array (a “tensor”) comprising neurons, timebins, and trials (fig. 3E, *left*). The data are then approximated by a three-way factorization, which is a straightforward generalization of the matrix factorization methods described in the last section. This produces three intertwined sets of low-dimensional factors (fig. 3E, *right*). Each triplet of factors describes a sub-population of neurons (blue), with a characteristic temporal firing pattern (red), with per-trial gain modulation (green).

When successful, tensor factorization can achieve a much more aggressive degree of dimensionality reduction

than matrix factorization. For example, one demonstration in Williams et al. [102<sup>•</sup>] on mouse prefrontal cortical dynamics shows that a tensor decomposition model utilizes 100-fold fewer parameters than a PCA model with nearly equal levels of reconstruction accuracy. Intuitively, this dramatic reduction in dimensionality (without sacrificing performance) facilitates visualization and interpretation of the data. For instance, Figure 3F shows one low-dimensional tensor component which identifies a sub-population of neurons whose activity grew in magnitude over the course of learning in a nonhuman primate learning to adapt to a visuomotor perturbation in a brain-computer interface task. By additionally reducing the dimensionality of the data across trials, tensor factorization is uniquely suited to pull out such trends in neural data. Like NMF, tensor factorizations are unique under weak assumptions (see [37]), particularly if nonnegativity constraints are additionally incorporated into the model. Thus, unlike PCA, tensor factorizations do not suffer from the “rotation problem” (see **Box 2** and fig. 2F), and so the extracted factors are more amenable to direct interpretation.

Finally, low-rank and non-negative matrix factorization, time-warping models, and tensor factorizations can all be seen as probabilistic models in which each trial is endowed with latent variables that specify its unique features. Hierarchical Bayesian models [24] generalize these notions by specifying a joint distribution over the entire dataset, combining information across trials to estimate global parameters (e.g., per-neuron factors) while simultaneously allowing latent variables (e.g., per-trial gain factors) to capture trial-by-trial variability. Of particular interest are hierarchical state space models [26, 42, 62], which model each trial as a sample of a stochastic dynamical system. For many problems of interest, we can formulate existing computational theories and hypotheses as dynamical systems models governing how a neural population’s activity evolves over time [41, 111]. Hierarchical state space models additionally generalize to include latent variables associated with each behavioral condition, brain region, recording session, subject, and so on. As neuroscience progresses to study more complex and naturalistic behaviors, hierarchical models that pool statistical power while capturing variability across trials, conditions, subjects, and sessions will be a critical component of our statistical toolkit.

## Open questions, challenges, and opportunities

What scientific discoveries might statistical models with single-trial resolution help us unlock? In general, we expect these methods to be most informative when behavioral performance is changing and unstable. Fluctuations in attention, which are thought to correlate with population-wide gain modulations in sensory areas, represent a relatively simple and statistically tractable example that we reviewed above in detail. More difficult tasks will often produce larger levels of trial-to-trial variability, reflecting changes in the animal’s uncertainty, strategy, and appraisal of evidence, all of which may evolve stochastically over time during internal deliberation [69]. Capturing signatures of these effects in neural dynamics remains challenging [8<sup>•</sup>], and simple modeling assumptions like gain modulation and time warping may be insufficient to capture the full scope of these complexities. State space models that optimize a set of stochastic differential equations governing neural dynamics are a promising and still evolving line of work on this subject [110, 111].

Trial-to-trial variability may also be heightened during incremental, long-term learning of complex tasks [15].

This, too, represents a promising application area in which statistical methods that capture the gradual emergence of learned dynamics could connect neurobiological data to classical theories of learning. For example, learning dynamics for hierarchically structured tasks (e.g. the categorization of objects into increasingly refined sub-categories) are expected to advance in discrete, step-like stages [50]—a prediction that was recently verified and theoretically characterized in artificial deep networks [77]. Similar opportunities originate in theories of reinforcement learning, which have heavily influenced neuroscience research for decades [56]. Here again, state space models offer a promising approach toward translating the constraints of competing theories into statistical models that can be tested against data [41].

More broadly, as illustrated in Figure 1A, we expect the field’s trend towards complex experiments with trial-limited regimes to increasingly necessitate the adoption of statistical models with single-trial resolution. Indeed, naturalistic animal behavior does not neatly sort itself into a discrete set of conditions with repeated trial structure. We can sometimes dispense with the notion of trials altogether and characterize neural activity by directly analyzing the streaming time series of neural and behavioral data—the discovery of place cells, grid cells, and head direction cells in the field of navigation are prominent examples of this [52]. However, the relationships between neural activity and behavior are not always so regular and predictable. In many cases, it is useful to extract *approximate trials* (i.e., inexact repetitions of a behavioral act or a pattern of neural population activity) from unstructured time series. Methods such as MoSeq [46], recurrent switching dynamical systems [43], dimensionality reduction and clustering based on wavelet features [5], and spike sequence detection models [45, 65, 67, 101], may be used for this purpose. However, the time series clustering problem these methods aim to solve is very challenging, and it is still an active area of research.

Finally, although descriptive statistical summaries are a key step towards understanding complex neural circuit dynamics, neuroscientists should aim to integrate trial-by-trial analyses more deeply into the design and execution of causal experiments. For example, the neural sub-populations identified by latent variable models could be differentially manipulated by emerging optogenetic stimulation protocols for large-scale populations [48], or targeted perturbations to brain-computer interfaces [57]. Such interventions will be critical if we hope to build causal links between neural population dynamics and animal behavior. Translating statistical modeling assumptions into biological terms—such as the connection between gain modulation and low-rank matrix factorization highlighted in this review—can help facilitate these important interactions between experimental and theoretical research.

## Acknowledgements

A.H.W. received funding support from the National Institutes of Health BRAIN initiative (1F32MH122998-01), and the Wu Tsai Stanford Neurosciences Institute Interdisciplinary Scholar Program. S.W.L. was supported by grants from the Simons Collaboration on the Global Brain (SCGB 697092) and the NIH BRAIN Initiative (U19NS113201 and R01NS113119).

## References

- [1] Afsheen Afshar, Gopal Santhanam, Byron M. Yu, Stephen I. Ryu, Maneesh Sahani, and Krishna V. Shenoy. “Single-Trial Neural Correlates of Arm Movement Preparation”. *Neuron* 71.3 (2011), pp. 555–564.
- [2] Bruno B. Averbeck, Peter E. Latham, and Alexandre Pouget. “Neural correlations, population coding and computation”. *Nature Reviews Neuroscience* 7.5 (2006), pp. 358–366.
- [3] Ramon Bartolo, Richard C. Saunders, Andrew R. Mitz, and Bruno B. Averbeck. “Information-Limiting Correlations in Large Neural Populations”. *Journal of Neuroscience* 40.8 (2020), pp. 1668–1678.
- [4] Brice Bathellier, Derek L. Buhl, Riccardo Accolla, and Alan Carleton. “Dynamic Ensemble Odor Coding in the Mammalian Olfactory Bulb: Sensory Information at Different Timescales”. *Neuron* 57.4 (2008), pp. 586–598.
- [5] Gordon J. Berman, Daniel M. Choi, William Bialek, and Joshua W. Shaevitz. “Mapping the stereotyped behaviour of freely moving fruit flies”. *Journal of The Royal Society Interface* 11.99 (2014), p. 20140672.
- [6] K. L. Briggman, H. D. I. Abarbanel, and W. B. Kristan. “Optical Imaging of Neuronal Populations During Decision-Making”. *Science* 307.5711 (2005), pp. 896–901.
- [7] Santiago A. Cadena, George H. Denfield, Edgar Y. Walker, Leon A. Gatys, Andreas S. Tolias, Matthias Bethge, and Alexander S. Ecker. “Deep convolutional models improve predictions of macaque V1 responses to natural images”. *PLOS Computational Biology* 15.4 (2019), pp. 1–27.
- [8•] Chandramouli Chandrasekaran, Joana Soldado-Magraner, Diogo Peixoto, William T. Newsome, Krishna V. Shenoy, and Maneesh Sahani. “Brittleness in model selection analysis of single neuron firing rates”. *bioRxiv* (2018).

The authors demonstrate that classical model selection techniques such as Akaike and Bayesian information criteria (AIC and BIC) can be surprisingly brittle and sensitive to model mismatch when analyzing single neuron recordings. To avoid these challenges, the authors argue that neuroscientists should apply and evaluate a broad variety of models, rather than a small number of pre-ordained hypotheses. Further, quantitative statistics that summarize model performance should be combined with data visualization and other qualitative measures of model agreement.
- [9] John K. Chapin, Karen A. Moxon, Ronald S. Markowitz, and Miguel A. L. Nicolelis. “Real-time control of a robot arm using simultaneously recorded neurons in the motor cortex”. *Nature Neuroscience* 2.7 (1999), pp. 664–670.
- [10] Mark M. Churchland, John P. Cunningham, Matthew T. Kaufman, Justin D. Foster, Paul Nuyujukian, Stephen I. Ryu, and Krishna V. Shenoy. “Neural population dynamics during reaching”. *Nature* 487.7405 (2012), pp. 51–56.

- [11] Mark M Churchland, M Yu Byron, Maneesh Sahani, and Krishna V Shenoy. “Techniques for extracting single-trial activity patterns from large-scale neural recordings”. *Current opinion in neurobiology* 17.5 (2007), pp. 609–618.
- [12•] Benjamin R. Cowley, Adam C. Snyder, Katerina Acar, Ryan C. Williamson, Byron M. Yu, and Matthew A. Smith. “Slow Drift of Neural Activity as a Signature of Impulsivity in Macaque Visual and Prefrontal Cortex”. *Neuron* 108.3 (2020), 551–567.e8.
- Cowley et al. show that behavioral measures of impulsivity drifted spontaneously and substantially over the course of several hours in monkeys performing a visual change detection task. This drift in performance was tightly correlated to drifts in neural firing rates in V4 and prefrontal cortex, which were independently identified by applying PCA on spike count residuals. This effect was not easily visible in single neurons, but could be reliably detected at the population-level. Altogether, these results are a powerful reminder that neurobiological systems and animal behaviors are often non-stationary, and demonstrate how trial-by-trial analyses can reveal a variety of additional details about the system.**
- [13] Brian M. Dekleva, Konrad P Kording, and Lee E. Miller. “Single reach plans in dorsal premotor cortex during a two-target task”. *Nature Communications* 9.1 (2018), p. 3556.
- [14] Ashesh K. Dhawale, Yohsuke R. Miyamoto, Maurice A. Smith, and Bence P. Ölveczky. “Adaptive Regulation of Motor Variability”. *Current Biology* 29.21 (2019), 3551–3562.e7.
- [15] Ashesh K. Dhawale, Maurice A. Smith, and Bence P. Ölveczky. “The Role of Variability in Motor Learning”. *Annual Review of Neuroscience* 40.1 (2017), pp. 479–498.
- [16] Ashesh K Dhawale, Rajesh Poddar, Steffen BE Wolff, Valentin A Normand, Evi Kopelowitz, and Bence P Ölveczky. “Automated long-term recording and analysis of neural activity in behaving animals”. *eLife* 6 (2017). Ed. by Andrew J King, e27702.
- [17] David Donoho and Victoria Stodden. “When Does Non-Negative Matrix Factorization Give a Correct Decomposition into Parts?” *Advances in Neural Information Processing Systems*. Ed. by S. Thrun, L. Saul, and B. Schölkopf. Vol. 16. MIT Press, 2004, pp. 1141–1148.
- [18] Yedidyah Dordek, Daniel Soudry, Ron Meir, and Dori Derdikman. “Extracting grid cell characteristics from place cell inputs using non-negative principal component analysis”. *eLife* 5 (2016). Ed. by Michael J Frank, e10094.
- [19] Lea Duncker and Maneesh Sahani. “Temporal alignment and latent Gaussian process factor inference in population spike trains”. *Advances in Neural Information Processing Systems* 31. Ed. by S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett. Curran Associates, Inc., 2018, pp. 10445–10455.

- [20] Gamaleldin F. Elsayed and John P. Cunningham. “Structure in neural population recordings: an expected byproduct of simpler phenomena?” *Nature Neuroscience* 20.9 (2017), pp. 1310–1318.
- [21] Katie A. Ferguson and Jessica A. Cardin. “Mechanisms underlying gain modulation in the cortex”. *Nature Reviews Neuroscience* 21.2 (2020), pp. 80–92.
- The authors provide an encyclopedic review of the experimental literature on gain modulation, which has a tight connection to the low-dimensional factor models of trial-to-trial variability. In visual cortex, where gain modulations are most comprehensively characterized, neuromodulatory inputs differentially target classes of GABAergic interneurons to shape synaptic integration in pyramidal cells and ultimately tune the gain of neural responses. There are possibly multiple forms of gain modulation, with different underlying cellular mechanisms, associated with locomotion and wakefulness/arousal. This potentially motivates the use of several low-rank factors (as done by [68•]) to model trial-to-trial variability.
- [22] Alfredo Fontanini and Donald B. Katz. “Behavioral States, Network States, and Sensory Response Variability”. *Journal of Neurophysiology* 100.3 (2008), pp. 1160–1168.
- [23] Jason P Gallivan, Craig S. Chapman, Daniel M. Wolpert, and J. Randall Flanagan. “Decision-making in sensorimotor control”. *Nature Reviews Neuroscience* 19.9 (2018), pp. 519–534.
- [24] Andrew Gelman, John B Carlin, Hal S Stern, David B Dunson, Aki Vehtari, and Donald B Rubin. *Bayesian Data Analysis, Third Edition*. en. CRC Press, 2013.
- [25] Nicolas Gillis. “The why and how of nonnegative matrix factorization”. *Regularization, Optimization, Kernels, and Support Vector Machines*. Ed. by Johan A.K. Suykens, Marco Signoretto, and Andreas Argyriou. Vol. 12. Machine Learning & Pattern Recognition Series. Chapman & Hall/CRC, 2014, pp. 257–291.
- [26] Joshua Glaser, Matthew Whiteway, John P Cunningham, Liam Paninski, and Scott Linderman. “Recurrent Switching Dynamical Systems Models for Multiple Interacting Neural Populations”. *Advances in Neural Information Processing Systems*. Ed. by H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin. Vol. 33. Curran Associates, Inc., 2020, pp. 14867–14878.
- [27] Joshua I. Gold and Michael N. Shadlen. “The Neural Basis of Decision Making”. *Annual Review of Neuroscience* 30.1 (2007), pp. 535–574.
- [28] Jorge Golowasch, Mark S. Goldman, L. F. Abbott, and Eve Marder. “Failure of Averaging in the Construction of a Conductance-Based Neuron Model”. *Journal of Neurophysiology* 87.2 (2002), pp. 1129–1131.
- [29] Nicholas G. Hatsopoulos, Catherine L. Ojakangas, Liam Paninski, and John P. Donoghue. “Information about movement direction obtained from synchronous activity of motor cortical neurons”. *Proceedings of the National Academy of Sciences* 95.26 (1998), pp. 15706–15711.

- [30] Jay Hegde and David C. Van Essen. “A Comparative Study of Shape Representation in Macaque Visual Areas V2 and V4”. *Cerebral Cortex* 17.5 (2006), pp. 1100–1116.
- [31] David I. Inouye, Eunho Yang, Genevera I. Allen, and Pradeep Ravikumar. “A review of multivariate distributions for count data derived from the Poisson distribution”. *WIREs Computational Statistics* 9.3 (2017), e1398.
- [32] Matthew T. Kaufman, Mark M. Churchland, and Krishna V. Shenoy. “The roles of monkey M1 neuron classes in movement preparation and execution”. *Journal of Neurophysiology* 110.4 (2013). PMID: 23699057, pp. 817–825.
- [33] Matthew T Kaufman, Mark M Churchland, Stephen I Ryu, and Krishna V Shenoy. “Vacillation, indecision and hesitation in moment-by-moment decoding of monkey motor cortex”. *eLife* 4 (2015). Ed. by Matteo Carandini, e04677.
- [34] Roozbeh Kiani, Christopher J. Cueva, John B. Reppas, and William T. Newsome. “Dynamics of Neural Population Responses in Prefrontal Cortex Indicate Changes of Mind on Single Trials”. *Current Biology* 24.13 (2014), pp. 1542–1547.
- [35] Dmitry Kobak, Wieland Brendel, Christos Constantinidis, Claudia E Feierstein, Adam Kepecs, Zachary F Mainen, Xue-Lian Qi, Ranulfo Romo, Naoshige Uchida, and Christian K Machens. “Demixed principal component analysis of neural population data”. *eLife* 5 (2016). Ed. by Mark CW van Rossum, e10989.
- [36] Adam Kohn, Ruben Coen-Cagli, Ingmar Kanitscheider, and Alexandre Pouget. “Correlations and Neuronal Population Information”. *Annual Review of Neuroscience* 39.1 (2016), pp. 237–256.
- [37] Tamara G. Kolda and Brett W. Bader. “Tensor Decompositions and Applications”. *SIAM Review* 51.3 (2009), pp. 455–500.
- [38] Daniel D. Lee and H. Sebastian Seung. “Learning the parts of objects by non-negative matrix factorization”. *Nature* 401.6755 (1999), pp. 788–791.
- [39] Joonyeol Lee, Mati Joshua, Javier F. Medina, and Stephen G. Lisberger. “Signal, Noise, and Variation in Neural and Sensory-Motor Latency”. *Neuron* 90.1 (2016), pp. 165–176.
- [40] Anthony Leonardo and Michale S. Fee. “Ensemble Coding of Vocal Control in Birdsong”. *Journal of Neuroscience* 25.3 (2005), pp. 652–661.
- [41] Scott W Linderman and Samuel J Gershman. “Using computational theory to constrain statistical models of neural data”. *Current Opinion in Neurobiology* 46 (2017). Computational Neuroscience, pp. 14–24.

- [42] Scott W Linderman, Annika L A Nichols, David M Blei, Manuel Zimmer, and Liam Paninski. “Hierarchical recurrent state space models reveal discrete and continuous dynamics of neural activity in *C. elegans*”. *bioRxiv* (2019).
- [43] Scott Linderman, Matthew Johnson, Andrew Miller, Ryan Adams, David Blei, and Liam Paninski. “Bayesian Learning and Inference in Recurrent Switching Linear Dynamical Systems”. Ed. by Aarti Singh and Jerry Zhu. Vol. 54. Proceedings of Machine Learning Research. Fort Lauderdale, FL, USA: PMLR, 2017, pp. 914–922.
- [44] Jian K. Liu, Helene M. Schreyer, Arno Onken, Fernando Rozenblit, Mohammad H. Khani, Vidhyasankar Krishnamoorthy, Stefano Panzeri, and Tim Gollisch. “Inference of neuronal functional circuitry with spike-triggered non-negative matrix factorization”. *Nature Communications* 8.1 (2017), p. 149.
- [45] Emily L Mackevicius, Andrew H Bahle, Alex H Williams, Shijie Gu, Natalia I Denisenko, Mark S Goldman, and Michale S Fee. “Unsupervised discovery of temporal sequences in high-dimensional datasets, with applications to neuroscience”. *eLife* 8 (2019). Ed. by Laura Colgin and Timothy E Behrens, e38471.
- [46] Jeffrey E. Markowitz, Winthrop F. Gillis, Celia C. Beron, Shay Q. Neufeld, Keiramarie Robertson, Neha D. Bhagat, Ralph E. Peterson, Emalee Peterson, Minsuk Hyun, Scott W. Linderman, Bernardo L. Sabatini, and Sandeep Robert Datta. “The Striatum Organizes 3D Behavior via Moment-to-Moment Action Selection”. *Cell* 174.1 (2018), 44–58.e17.
- [47] Jesse D. Marshall, Diego E. Aldarondo, Timothy W. Dunn, William L. Wang, Gordon J. Berman, and Bence P. Ölveczky. “Continuous Whole-Body 3D Kinematic Recordings across the Rodent Behavioral Repertoire”. *Neuron* () .
- [48] James H. Marshel, Yoon Seok Kim, Timothy A. Machado, Sean Quirin, Brandon Benson, Jonathan Kadmon, Cephra Raja, Adelaida Chibukhchyan, Charu Ramakrishnan, Masatoshi Inoue, Janelle C. Shane, Douglas J. McKnight, Susumu Yoshizawa, Hideaki E. Kato, Surya Ganguli, and Karl Deisseroth. “Cortical layer-specific critical dynamics triggering perception”. *Science* 365.6453 (2019).
- [49] Paul Masset, Torben Ott, Armin Lak, Junya Hirokawa, and Adam Kepcs. “Behavior- and Modality-General Representation of Confidence in Orbitofrontal Cortex”. *Cell* 182.1 (2020), 112–126.e18.
- [50] J. L. McClelland. “A connectionist perspective on knowledge and development.” Developing cognitive competence: New approaches to process modeling. Hillsdale, NJ, US: Lawrence Erlbaum Associates, Inc, 1995, pp. 157–204.
- [51] Gal Mishne, Ronen Talmon, Ron Meir, Jackie Schiller, Maria Lavzin, Uri Dubin, and Ronald R Coifman. “Hierarchical coupled-geometry analysis for neuronal structure and activity pattern discovery”. *IEEE Journal of Selected Topics in Signal Processing* 10.7 (2016), pp. 1238–1253.

- [52] Edvard I. Moser, Emilio Kropff, and May-Britt Moser. “Place Cells, Grid Cells, and the Brain’s Spatial Representation System”. *Annual Review of Neuroscience* 31.1 (2008), pp. 69–89.
- [53] Simon Musall, Matthew T. Kaufman, Ashley L. Juavinett, Steven Gluf, and Anne K. Churchland. “Single-trial neural dynamics are dominated by richly varied movements”. *Nature Neuroscience* 22.10 (2019), pp. 1677–1686.
- [54] A. M. Ni, D. A. Ruff, J. J. Alberts, J. Symmonds, and M. R. Cohen. “Learning and attention reveal a general relationship between population activity and behavior”. *Science* 359.6374 (2018), pp. 463–465.
- [55] Christopher M. Niell and Michael P. Stryker. “Modulation of Visual Responses by Behavioral State in Mouse Visual Cortex”. *Neuron* 65.4 (2010), pp. 472–479.
- [56] Yael Niv. “Reinforcement learning in the brain”. *Journal of Mathematical Psychology* 53.3 (2009). Special Issue: Dynamic Decision Making, pp. 139–154.
- [57] Emily R. Oby, Matthew D. Golub, Jay A. Hennig, Alan D. Degenhart, Elizabeth C. Tyler-Kabara, Byron M. Yu, Steven M. Chase, and Aaron P. Batista. “New neural activity patterns emerge with long-term learning”. *Proceedings of the National Academy of Sciences* 116.30 (2019), pp. 15210–15215.
- [58] Kenichi Ohki, Sooyoung Chung, Yeang H. Ch’ng, Prakash Kara, and R. Clay Reid. “Functional imaging with cellular resolution reveals precise micro-architecture in visual cortex”. *Nature* 433.7026 (2005), pp. 597–603.
- [59] Arno Onken, Jian K. Liu, P. P. Chamanthi R. Karunasekara, Ioannis Delis, Tim Gollisch, and Stefano Panzeri. “Using Matrix and Tensor Factorizations for the Single-Trial Analysis of Population Spike Trains”. *PLOS Computational Biology* 12.11 (2016), pp. 1–46.
- [60] Chethan Pandarinath, Daniel J. O’Shea, Jasmine Collins, Rafal Jozefowicz, Sergey D. Stavisky, Jonathan C. Kao, Eric M. Trautmann, Matthew T. Kaufman, Stephen I. Ryu, Leigh R. Hochberg, Jaimie M. Henderson, Krishna V. Shenoy, L. F. Abbott, and David Sussillo. “Inferring single-trial neural population dynamics using sequential auto-encoders”. *Nature Methods* 15.10 (2018), pp. 805–815.
- [61] Liam Paninski. “Maximum likelihood estimation of cascade point-process neural encoding models”. *Network: Computation in Neural Systems* 15.4 (2004), pp. 243–262.
- [62] Liam Paninski, Yashar Ahmadian, Daniel Gil Ferreira, Shinsuke Koyama, Kamiar Rahnama Rad, Michael Vidne, Joshua Vogelstein, and Wei Wu. “A new look at state-space models for neural data”. *Journal of Computational Neuroscience* 29.1 (2010), pp. 107–126.
- [63] Junchol Park, Luke T. Coddington, and Joshua T. Dudman. “Basal Ganglia Circuits for Action Specification”. *Annual Review of Neuroscience* 43.1 (2020). PMID: 32303147, pp. 485–507.

- [64] Stan L. Pashkovski, Giuliano Iurilli, David Brann, Daniel Chicharro, Kristen Drumme, Kevin M. Franks, Stefano Panzeri, and Sandeep Robert Datta. “Structure and flexibility in cortical representations of odour space”. *Nature* 583.7815 (2020), pp. 253–258.
- [65] Sven Peter, Elke Kirschbaum, Martin Both, Lee Campbell, Brandon Harvey, Conor Heins, Daniel Durstewitz, Ferran Diego, and Fred A Hamprecht. “Sparse convolutional coding for neuronal assembly detection”. *Advances in Neural Information Processing Systems 30*. Ed. by I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett. Curran Associates, Inc., 2017, pp. 3675–3685.
- [66] Eftychios A. Pnevmatikakis, Daniel Soudry, Yuanjun Gao, Timothy A. Machado, Josh Merel, David Pfau, Thomas Reardon, Yu Mu, Clay Lacefield, Weijian Yang, Misha Ahrens, Randy Bruno, Thomas M. Jessell, Darcy S. Peterka, Rafael Yuste, and Liam Paninski. “Simultaneous Denoising, Deconvolution, and Demixing of Calcium Imaging Data”. *Neuron* 89.2 (2016), pp. 285–299.
- [67] Pietro Quaglio, Vahid Rostami, Emiliano Torre, and Sonja Grün. “Methods for identification of spike patterns in massively parallel spike trains”. *Biological Cybernetics* 112.1 (2018), pp. 57–80.
- [68•] Neil C Rabinowitz, Robbe L Goris, Marlene Cohen, and Eero P Simoncelli. “Attention stabilizes the shared gain of V4 populations”. *eLife* 4 (2015). Ed. by Matteo Carandini, e08998.  
**The authors develop a low-dimensional factor model to capture how attentional changes modulate neural firing rates in area V4 of visual cortex. The model shares many features with other highlighted work—it includes a term to capture slow drift [12•], and exploits the tensor structure of the dataset [102•]. The low-dimensional factors extracted by this model improve predictive log-likelihood over models without trial-by-trial modulations. Further, the extracted factors are interpretable and relate concrete changes in neural processing to attentional cues, rewards, and choices.**
- [69] Arbora Resulaj, Roozbeh Kiani, Daniel M. Wolpert, and Michael N. Shadlen. “Changes of mind in decision-making”. *Nature* 461.7261 (2009), pp. 263–266.
- [70] John H. Reynolds and David J. Heeger. “The Normalization Model of Attention”. *Neuron* 61.2 (2009), pp. 168–185.
- [71] Rajeev V. Rikhye and Mriganka Sur. “Spatial Correlations in Natural Scenes Modulate Response Reliability in Mouse Visual Cortex”. *Journal of Neuroscience* 35.43 (2015), pp. 14661–14680.
- [72] Nicholas A. Roy, Ji Hyun Bak, The International Brain Laboratory, Athena Akrami, Carlos D. Brody, and Jonathan W. Pillow. “Extracting the Dynamics of Behavior in Decision-Making Experiments”. *bioRxiv* (2020).
- [73] Douglas A. Ruff and Marlene R. Cohen. “Global Cognitive Factors Modulate Correlated Response Variability between V4 Neurons”. *Journal of Neuroscience* 34.49 (2014), pp. 16408–16416.

- [74•] Oleg I. Rumyantsev, Jérôme A. Lecoq, Oscar Hernandez, Yanping Zhang, Joan Savall, Radosław Chrapkiewicz, Jane Li, Hongkui Zeng, Surya Ganguli, and Mark J. Schnitzer. “Fundamental bounds on the fidelity of sensory cortical coding”. *Nature* 580.7801 (2020), pp. 100–105.
- By simultaneously recording from very large neural populations, the authors demonstrate the negative impact of correlated noise on single-trial decoding of oriented gratings in visual cortex. This provides direct experimental insights into a longstanding topic of interest. In addition, the authors provide an insightful theoretical analysis—under certain assumptions, they show that as more neurons are simultaneously recorded, fewer trials are needed to approximate the important eigenvalues of the noise covariance. See the contemporaneous study by Bartolo et al. [3], for a similar result in a different animal model and behavioral task.
- [75] Abigail A. Russo, Sean R. Bittner, Sean M. Perkins, Jeffrey S. Seely, Brian M. London, Antonio H. Lara, Andrew Miri, Najja J. Marshall, Adam Kohn, Thomas M. Jessell, Laurence F. Abbott, John P. Cunningham, and Mark M. Churchland. “Motor Cortex Embeds Muscle-like Commands in an Untangled Population Response”. *Neuron* 97.4 (2018), 953–966.e8.
- [76] Abigail A. Russo, Ramin Khajeh, Sean R. Bittner, Sean M. Perkins, John P. Cunningham, L. F. Abbott, and Mark M. Churchland. “Neural Trajectories in the Supplementary Motor Area and Motor Cortex Exhibit Distinct Geometries, Compatible with Different Classes of Computation”. *Neuron* 107.4 (2020), 745–758.e6.
- [77] Andrew M. Saxe, James L. McClelland, and Surya Ganguli. “A mathematical theory of semantic development in deep neural networks”. *Proceedings of the National Academy of Sciences* 116.23 (2019), pp. 11537–11546.
- [78] Shreya Saxena, Ian Kinsella, Simon Musall, Sharon H. Kim, Jozsef Meszaros, David N. Thibodeaux, Carla Kim, John Cunningham, Elizabeth M. C. Hillman, Anne Churchland, and Liam Paninski. “Localized semi-nonnegative matrix factorization (LocaNMF) of widefield calcium imaging data”. *PLOS Computational Biology* 16.4 (2020), pp. 1–28.
- [79] Odelia Schwartz, Jonathan W Pillow, Nicole C Rust, and Eero P Simoncelli. “Spike-triggered neural characterization”. *Journal of vision* 6.4 (2006), pp. 13–13.
- [80] Jeffrey S. Seely, Matthew T. Kaufman, Stephen I. Ryu, Krishna V. Shenoy, John P. Cunningham, and Mark M. Churchland. “Tensor Analysis Reveals Distinct Population Structure that Parallels the Different Computational Roles of Areas M1 and V1”. *PLOS Computational Biology* 12.11 (2016), pp. 1–34.
- [81•] Nishal P Shah, Nora Brackbill, Colleen Rhoades, Alexandra Kling, Georges Goetz, Alan M Litke, Alexander Sher, Eero P Simoncelli, and EJ Chichilnisky. “Inference of nonlinear receptive field subunits with spike-triggered clustering”. *eLife* 9 (2020). Ed. by Tatyana O Sharpee and Joshua I Gold, e45743. The authors extend the spike-triggered NMF model developed by Liu et al. [44] into a fully probabilistic clustering model with additional regularization terms and spiking nonlinearities.

Since the cluster assignments are attributed probabilistically, the model retains a “soft clustering” interpretation, which is similar to NMF. When applied to data from retinal ganglion cells, the model partitions the receptive field into functional subunits that likely reflect bipolar cell inputs. When the model is fit to neighboring ganglion cells, subunits are discovered in consistent locations, presumably reflecting shared pre-synaptic inputs. Overall, this demonstrates the ability of nonnegative factor models to identify interpretable biological structure, purely from neural activity.

- [82•] Hansem Sohn, Devika Narain, Nicolas Meirhaeghe, and Mehrdad Jazayeri. “Bayesian Computation through Cortical Latent Dynamics”. *Neuron* 103.5 (2019), 934–947.e5.  
**The authors investigate how prior beliefs and sensory evidence are integrated into neural representations in prefrontal cortex during a sensorimotor time keeping task. They argue that the amount of curvature in neural firing rate trajectories can affect this integration. Intriguingly, and relevant to this review, the authors show evidence that trial-to-trial variability is preferentially aligned with vectors tangent to the path of the trial-averaged trajectory. This geometric orientation of the variability is consistent with a Bayesian model—uncertainty in the response (neural variability in the output-potent dimensions) decreases as the magnitude of accumulated evidence (position along the curved neural manifold) increases.**
- [83] Ben Sorscher, Gabriel Mel, Surya Ganguli, and Samuel Ocko. “A unified theory for the origin of grid cells through the lens of pattern formation”. *Advances in Neural Information Processing Systems*. Ed. by H. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and R. Garnett. Vol. 32. Curran Associates, Inc., 2019, pp. 10003–10013.
- [84] H Spitzer, R Desimone, and J Moran. “Increased attention enhances both behavioral and neuronal performance”. *Science* 240.4850 (1988), pp. 338–340.
- [85] Nicholas A. Steinmetz, Cagatay Aydin, Anna Lebedeva, Michael Okun, Marius Pachitariu, Marius Bauza, Maxime Beau, Jai Bhagat, Claudia Böhm, Martijn Broux, Susu Chen, Jennifer Colonell, Richard J. Gardner, Bill Karsh, Dimitar Kostadinov, Carolina Mora-Lopez, Junchol Park, Jan Putzeys, Britton Sauerbrei, Rik J. J. van Daal, Abraham Z. Vollan, Marleen Welkenhuysen, Zhiwen Ye, Joshua Dudman, Barundeb Dutta, Adam W. Hantman, Kenneth D. Harris, Albert K. Lee, Edvard I. Moser, John O’Keefe, Alfonso Renart, Karel Svoboda, Michael Häusser, Sebastian Haesler, Matteo Carandini, and Timothy D. Harris. “Neuropixels 2.0: A miniaturized high-density probe for stable, long-term brain recordings”. *bioRxiv* (2020).
- [86] Nicholas A. Steinmetz, Peter Zatka-Haas, Matteo Carandini, and Kenneth D. Harris. “Distributed coding of choice, action and engagement across the mouse brain”. *Nature* 576.7786 (2019), pp. 266–273.
- [87] Carsen Stringer, Marius Pachitariu, Nicholas Steinmetz, Matteo Carandini, and Kenneth D. Harris. “High-dimensional geometry of population responses in visual cortex”. *Nature* 571.7765 (2019), pp. 361–365.

- [88] Carsen Stringer, Marius Pachitariu, Nicholas Steinmetz, Charu Bai Reddy, Matteo Carandini, and Kenneth D. Harris. “Spontaneous behaviors drive multidimensional, brainwide activity”. *Science* 364.6437 (2019).
- [89] Dawn M. Taylor, Stephen I. Helms Tillery, and Andrew B. Schwartz. “Direct Cortical Control of 3D Neuroprosthetic Devices”. *Science* 296.5574 (2002), pp. 1829–1832.
- [90•] Madeleine Udell, Corinne Horn, Reza Zadeh, and Stephen Boyd. “Generalized Low Rank Models”. *Foundations and Trends in Machine Learning* 9.1 (2016), pp. 1–118.  
**The authors synthesize a large number unsupervised learning models—PCA, NMF, k-means clustering, and others—into a common matrix factorization paradigm, and use this foundation to develop a variety of model extensions. This didactic monograph is a useful entry point to neuroscientists who are interested in navigating the plethora of models that could be exploited for neural population analysis.**
- [91] Justus V. Verhagen, Daniel W. Wesson, Theoden I. Netoff, John A. White, and Matt Wachowiak. “Sniffing controls an adaptive filter of sensory input to the olfactory bulb”. *Nature Neuroscience* 10.5 (2007), pp. 631–639.
- [92•] Roman Vershynin. *High-dimensional probability: An introduction with applications in data science*. Vol. 47. Cambridge university press, 2018.  
**This book packs an extensive introduction to high-dimensional statistics into fewer than 300 pages, while also providing a highly approachable and an engaging read. We believe this is a must-read for theoretical neuroscientists with an interest in the statistical foundations of single-trial analysis. We also recommend Martin Wainwright’s related book [97], which covers many additional topics and applications.**
- [93] Roman Vershynin. “How Close is the Sample Covariance Matrix to the Actual Covariance Matrix?” *Journal of Theoretical Probability* 25.3 (2012), pp. 655–686.
- [94] Martin Vinck, Renata Batista-Brito, Ulf Knoblich, and Jessica A. Cardin. “Arousal and Locomotion Make Distinct Contributions to Cortical Activity Patterns and Visual Encoding”. *Neuron* 86.3 (2015), pp. 740–754.
- [95] Saskia E. J. de Vries, Jerome A. Lecoq, Michael A. Buice, Peter A. Groblewski, Gabriel K. Ocker, Michael Oliver, David Feng, Nicholas Cain, Peter Ledochowitsch, Daniel Millman, Kate Roll, Marina Garrett, Tom Keenan, Leonard Kuan, Stefan Mihalas, Shawn Olsen, Carol Thompson, Wayne Wakeman, Jack Waters, Derric Williams, Chris Barber, Nathan Berbesque, Brandon Blanchard, Nicholas Bowles, Shiella D. Caldejon, Linzy Casal, Andrew Cho, Sissy Cross, Chinh Dang, Tim Dolbeare, Melise Edwards, John Galbraith, Nathalie Gaudreault, Terri L. Gilbert, Fiona Griffin, Perry Hargrave, Robert Howard, Lawrence Huang, Sean Jewell, Nika Keller, Ulf Knoblich, Josh D. Larkin, Rachael Larsen, Chris Lau, Eric Lee, Felix Lee, Arielle Leon, Lu Li, Fuhui Long, Jennifer Luviano, Kyla Mace, Thuyanh Nguyen,

- Jed Perkins, Miranda Robertson, Sam Seid, Eric Shea-Brown, Jianghong Shi, Nathan Sjoquist, Cliff Slaughterbeck, David Sullivan, Ryan Valenza, Casey White, Ali Williford, Daniela M. Witten, Jun Zhuang, Hongkui Zeng, Colin Farrell, Lydia Ng, Amy Bernard, John W. Phillips, R. Clay Reid, and Christof Koch. “A large-scale standardized physiological survey reveals functional organization of the mouse visual cortex”. *Nature Neuroscience* 23.1 (2020), pp. 138–151.
- [96] Saurabh Vyas, Matthew D. Golub, David Sussillo, and Krishna V. Shenoy. “Computation Through Neural Population Dynamics”. *Annual Review of Neuroscience* 43.1 (2020), pp. 249–275.
  - [97] Martin J. Wainwright. *High-Dimensional Statistics: A Non-Asymptotic Viewpoint*. Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge University Press, 2019.
  - [98] M. R. Whiteway, K. Socha, V. Bonin, and D. A. Butts. “Characterizing the nonlinear structure of shared variability in cortical neuron populations using latent variable models”. *Neuron Behav Data Anal Theory* 3.1 (2019).
  - [99] Matthew R. Whiteway and Daniel A. Butts. “Revealing unobserved factors underlying cortical activity with a rectified latent variable model applied to neural population recordings”. *Journal of Neurophysiology* 117.3 (2017). PMID: 27927786, pp. 919–936.
  - [100] Alex H. Williams. “Combining tensor decomposition and time warping models for multi-neuronal spike train analysis”. *bioRxiv* (2020).
  - [101] Alex H Williams, Anthony Degleris, Yixin Wang, and Scott W Linderman. “Point process models for sequence detection in high-dimensional neural spike trains”. *arXiv preprint arXiv:2010.04875* (2020).
  - [102•] Alex H Williams, Tony Hyun Kim, Forea Wang, Saurabh Vyas, Stephen I Ryu, Krishna V Shenoy, Mark Schnitzer, Tamara G Kolda, and Surya Ganguli. “Unsupervised discovery of demixed, low-dimensional neural dynamics across multiple timescales through tensor component analysis”. *Neuron* 98.6 (2018), pp. 1099–1115.
- Describes tensor decomposition (“tensor components analysis”) as a general-purpose, unsupervised method for extracting within-trial and across-trial components of neural population activity. Diverse applications are demonstrated in data from artificial neural networks, calcium imaging data from rodent prefrontal cortex, and primate motor cortex. See Kolda & Bader [37] for a canonical review and introduction to tensor decompositions.**
- [103•] Alex H Williams, Ben Poole, Niru Maheswaranathan, Ashesh K Dhawale, Tucker Fisher, Christopher D Wilson, David H Brann, Eric M Trautmann, Stephen Ryu, Roman Shusterman, et al. “Discovering precise temporal patterns in large-scale neural recordings through robust and interpretable time warping”. *Neuron* 105.2 (2020), pp. 246–259.
- Describes a flexible framework for aligning neural dynamics across trials by time warping, and demonstrates its effectiveness on several datasets spanning sensory and motor systems from**

rodents and primates. See Duncker et al. [19] for an alternative approach, and Williams [100] for notes on integrating time warping into tensor decomposition models of multi-trial data.

- [104] Christopher D. Wilson, Gabriela O. Serrano, Alexei A. Koulakov, and Dmitry Rinberg. “A primacy code for odor identity”. *Nature Communications* 8.1 (2017), p. 1477.
- [105<sup>\*</sup>] Anqi Wu, Stan Pashkovski, Sandeep R Datta, and Jonathan W Pillow. “Learning a latent manifold of odor representations from neural responses in piriform cortex”. *Advances in Neural Information Processing Systems*. Ed. by S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett. Vol. 31. Curran Associates, Inc., 2018, pp. 5378–5388.  
The authors model the responses of  $N = 500$  neurons in piriform cortex to  $P = 66$  odorant stimuli using Gaussian Process latent variable models. To simultaneously model fluctuations across odors and neurons on each batch of trials, they decompose the overall  $NP \times NP$  covariance matrix as  $\Sigma \approx \Sigma_1 \otimes \Sigma_2$ , where  $\Sigma_1$  is an  $N \times N$  covariance matrix describing correlations between neurons, and  $\Sigma_2$  is a  $P \times P$  covariance matrix describing correlations between odors. This substantially reduces the number of parameters and is an instructive example that could be repeated in many other settings for multi-modal neural data analysis.
- [106] Byron M. Yu, John P. Cunningham, Gopal Santhanam, Stephen I. Ryu, Krishna V. Shenoy, and Maneesh Sahani. “Gaussian-Process Factor Analysis for Low-Dimensional Single-Trial Analysis of Neural Population Activity”. *Journal of Neurophysiology* 102.1 (2009), pp. 614–635.
- [107] Pengcheng Zhou, Shanna L Resendez, Jose Rodriguez-Romaguera, Jessica C Jimenez, Shay Q Neufeld, Andrea Giovannucci, Johannes Friedrich, Eftychios A Pnevmatikakis, Garret D Stuber, Rene Hen, Mazen A Kheirbek, Bernardo L Sabatini, Robert E Kass, and Liam Paninski. “Efficient and accurate extraction of in vivo calcium signals from microendoscopic video data”. *eLife* 7 (2018). Ed. by David C Van Essen, e28728.
- [108<sup>\*</sup>] Marinka Zitnik, Francis Nguyen, Bo Wang, Jure Leskovec, Anna Goldenberg, and Michael M Hoffman. “Machine learning for integrating data in biology and medicine: Principles, practice, and opportunities”. *Information Fusion* 50 (2019), pp. 71–91.  
Neural dynamics vary in systematic ways across trials, across timebins within a trial, and across conditions. It may soon become common to collect additional data—e.g., single-cell RNA sequencing or connectivity information—alongside neural activity measurements. As we advance in this direction, it will become increasingly important to build unified models that pool shared statistical information across these multiple data source. To tackle this challenge, neurostatisticians may find inspiration in neighboring fields, such as computational genomics and health care, who have been grappling with similar problems. This review by Zitnik et al. provides a useful entry point to this literature on “data fusion” and “multimodal learning.”

- [109] Ehud Zohary, Michael N. Shadlen, and William T. Newsome. “Correlated neuronal discharge rate and its implications for psychophysical performance”. *Nature* 370.6485 (1994), pp. 140–143.
- [110] David M Zoltowski, Kenneth W Latimer, Jacob L Yates, Alexander C Huk, and Jonathan W Pillow. “Discrete stepping and nonlinear ramping dynamics underlie spiking responses of LIP neurons during decision-making”. *Neuron* 102.6 (2019), pp. 1249–1258.
- [111] David M Zoltowski, Jonathan W Pillow, and Scott W Linderman. “A general recurrent state space framework for modeling neural dynamics during decision-making”. *Proceedings of the 37th International Conference on Machine Learning*. Ed. by Hal DaumÃ© III and Aarti Singh. Vol. 119. Proceedings of Machine Learning Research. Virtual: PMLR, 2020, pp. 11680–11691.
- [112] Hui Zou, Trevor Hastie, and Robert Tibshirani. “Sparse Principal Component Analysis”. *Journal of Computational and Graphical Statistics* 15.2 (2006), pp. 265–286.